

Group member name(s): Ahmed Arian Sajid, Sabrina Sultana, Sharmin Ahmad
Group member UID(s): 235061, 235062, 230239

Advanced Statistical Learning

Answer to Exercise 1.a

Consider the hypothesis spaces defined by the rule

$$\mathcal{H}_i := \{f(x) = \sum_{k=0}^i \theta_k x^k \mid \theta \in \mathbb{R}^{i+1}\} \quad (1)$$

For fitting model using the hypothesis spaces are,

$$H_1 = f(x) = \theta_0 + \theta_1 x \quad (2)$$

$$H_2 = f(x) = \theta_0 + \theta_1 x + \theta_2 x^2 \quad (3)$$

$$H_3 = f(x) = \theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 \quad (4)$$

$$H_5 = f(x) = \theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4 + \theta_5 x^5 \quad (5)$$

$$H_8 = f(x) = \theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4 + \theta_5 x^5 + \theta_6 x^6 + \theta_7 x^7 + \theta_8 x^8 \quad (6)$$

$$H_{12} = f(x) = \theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4 + \theta_5 x^5 + \theta_6 x^6 + \theta_7 x^7 + \theta_8 x^8 + \theta_9 x^9 + \theta_{10} x^{10} + \theta_{11} x^{11} + \theta_{12} x^{12} \quad (7)$$

visualization of the data and fitted models below.

```
mydt <- read.csv("C:/Users/Sabrina/Downloads/fitting.csv")

hyp1 <- lm(y ~ x, data = mydt)
hyp2 <- lm(y ~ x + I(x^2), data = mydt)
hyp3 <- lm(y ~ x + I(x^2) + I(x^3), data = mydt)
hyp5 <- lm(y ~ x + I(x^2) + I(x^3) + I(x^4) + I(x^5), data = mydt)
hyp8 <- lm(y ~ x + I(x^2) + I(x^3) + I(x^4) + I(x^5) + I(x^6) + I(x^7)
+ I(x^8), data = mydt)
hyp12 <- lm(y ~ x + I(x^2) + I(x^3) + I(x^4) + I(x^5) + I(x^6) + I(x^7)
+ I(x^8) + I(x^9) + I(x^10) + I(x^11) + I(x^12), data = mydt)

fn1 <- function(x) coef(hyp1)[1] + coef(hyp1)[2] * x
fn2 <- function(x) coef(hyp2)[1] + coef(hyp2)[2] * x + coef(hyp2)[3] * x^2
fn3 <- function(x) coef(hyp3)[1] + coef(hyp3)[2] * x + coef(hyp3)[3] * x^2 +
coef(hyp3)[4] * x^3
fn5 <- function(x) coef(hyp5)[1] + coef(hyp5)[2] * x + coef(hyp5)[3] * x^2 +
coef(hyp5)[4] * x^3 + coef(hyp5)[5] * x^4 + coef(hyp5)[6] * x^5
fn8 <- function(x) coef(hyp8)[1] + coef(hyp8)[2] * x + coef(hyp8)[3] * x^2 +
coef(hyp8)[4] * x^3 + coef(hyp8)[5] * x^4 + coef(hyp8)[6] * x^5 +
coef(hyp8)[7] * x^6 + coef(hyp8)[8] * x^7 + coef(hyp8)[9] * x^8
fn12 <- function(x) coef(hyp12)[1] + coef(hyp12)[2] * x + coef(hyp12)[3] * x^2 +
coef(hyp12)[4] * x^3 + coef(hyp12)[5] * x^4 + coef(hyp12)[6] * x^5 +
coef(hyp12)[7] * x^6 + coef(hyp12)[8] * x^7 + coef(hyp12)[9] * x^8 +
coef(hyp12)[10]*x^9 + coef(hyp12)[11]*x^10 + coef(hyp12)[12]* x^11 +
coef(hyp12)[13]*x^12

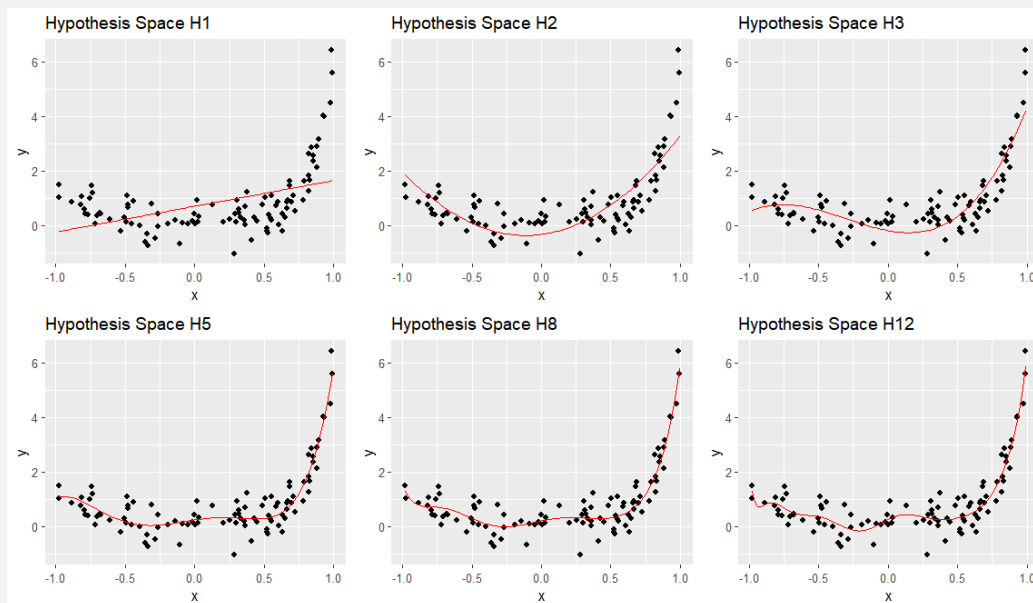
library(ggplot2)
```

```

val <- ggplot(aes(y = y, x = x), data = mydt) + geom_point()

h1 <- val + geom_function(fun = fn1, col = "red") + ggtitle("Hypothesis Space H1")
h2 <- val + geom_function(fun = fn2, col = "red") + ggtitle("Hypothesis Space H2")
h3 <- val + geom_function(fun = fn3, col = "red") + ggtitle("Hypothesis Space H3")
h5 <- val + geom_function(fun = fn5, col = "red") + ggtitle("Hypothesis Space H5")
h8 <- val + geom_function(fun = fn8, col = "red") + ggtitle("Hypothesis Space H8")
h12 <- val + geom_function(fun = fn12, col = "red") + ggtitle("Hypothesis Space H12")
install.packages("gridExtra")
library(gridExtra)
grid.arrange(h1,h2,h3,h5,h8,h12, nrow = 2)

```



Answer to Exercise 1.b

As we can see, there are 6 hypothesis spaces. Among them H_1 is linear regression line and H_2 is quadratic function line. These two lines hardly touch the data points. These two hypothesis spaces are under fitting.

H_3 is better and but this also can be said under fitting.

H_5 is considered to be the best model. It represents all the dataset nicely.

H_{12} and H_8 show 'wigglyness' in the data set. So these two hypothesis spaces are considered as over fitted model.

Answer to Exercise 1.c

Here we first take the data set and fit the models ($H_1, H_2, H_3, H_5, H_8, H_{12}$) and split the data into train dataset and test dataset. Then calculate the error (MSE)

```

mydt <- read.csv("C:/Users/Sabrina/Downloads/test.csv")

train <- mydt[1:80,]
test <- mydt[81:100,]

hyp1 <- lm(y ~ x, data = train)
hyp2 <- lm(y ~ x + I(x^2), data = train)
hyp3 <- lm(y ~ x + I(x^2) + I(x^3), data = train)

```

```

hyp5 <- lm(y ~ x + I(x^2) + I(x^3) + I(x^4) + I(x^5), data = train)
hyp8 <- lm(y ~ x + I(x^2) + I(x^3) + I(x^4) + I(x^5) + I(x^6) +
I(x^7) + I(x^8), data = train)
hyp12 <-lm(y ~ x + I(x^2) + I(x^3) + I(x^4) + I(x^5) + I(x^6) +
I(x^7) + I(x^8) + I(x^9) + I(x^10) + I(x^11) + I(x^12), data = train)

fits<- list(hyp1, hyp2, hyp3, hyp5, hyp8, hyp12)
MSE <- lapply(fits, function(fit) sum((predict(fit, newdata = test) - test$y)^2))
MSE <- data.frame(MSE = unlist(MSE), models = c("H1", "H2", "H3", "H5", "H8", "H12"))
MSE

```

MSE models		
1	25.897546	H1
2	9.410835	H2
3	6.807164	H3
4	4.041775	H5
5	9.117652	H8
6	23.569022	H12

Here, we can see that H_5 gives the less error than the other model.

Answer to Exercise 2

- (1 point) Plot the true decision boundary for the data situation.

Solution:

```

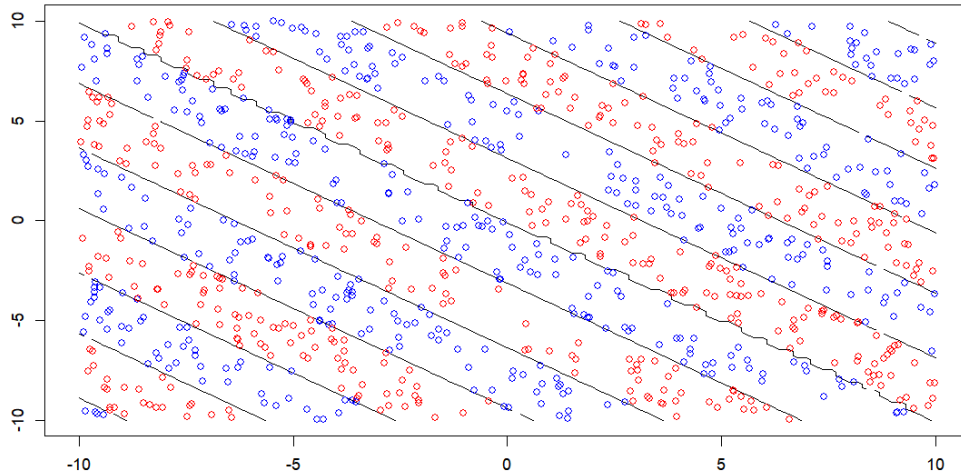
n <- 1000
x1 <- runif(n, -10, 10)
x2 <- runif(n, -10, 10)
y <- sin(x1 + x2) < 0

x1_grid <- seq(-10, 10, length.out = 100)
x2_grid <- seq(-10, 10, length.out = 100)
grid <- expand.grid(x1_grid, x2_grid)

z <- apply(grid, 1, function(p) as.numeric(sin(p[1] + p[2]) < 0))

contour(x1_grid, x2_grid, matrix(z, nrow = length(x2_grid),
ncol = length(x1_grid)), levels = 0.5, labels = "", xlim = c(-10, 10),
ylim = c(-10, 10))
points(x1[y], x2[y], col = "blue")
points(x1[!y], x2[!y], col = "red")

```



- (1 point) Create an explicit dataset with $n = 250$. Visualize the decision boundary for $k = 1, 5, 10, 100$. What do you observe?

Solution:

```
n <- 250
x1 <- runif(n, -10, 10)
x2 <- runif(n, -10, 10)
y <- as.numeric(sin(x1 + x2) < 0)
data <- data.frame(x1, x2, y)

x1_range <- seq(min(x1), max(x1), length.out = 100)
x2_range <- seq(min(x2), max(x2), length.out = 100)
grid <- expand.grid(x1 = x1_range, x2 = x2_range)

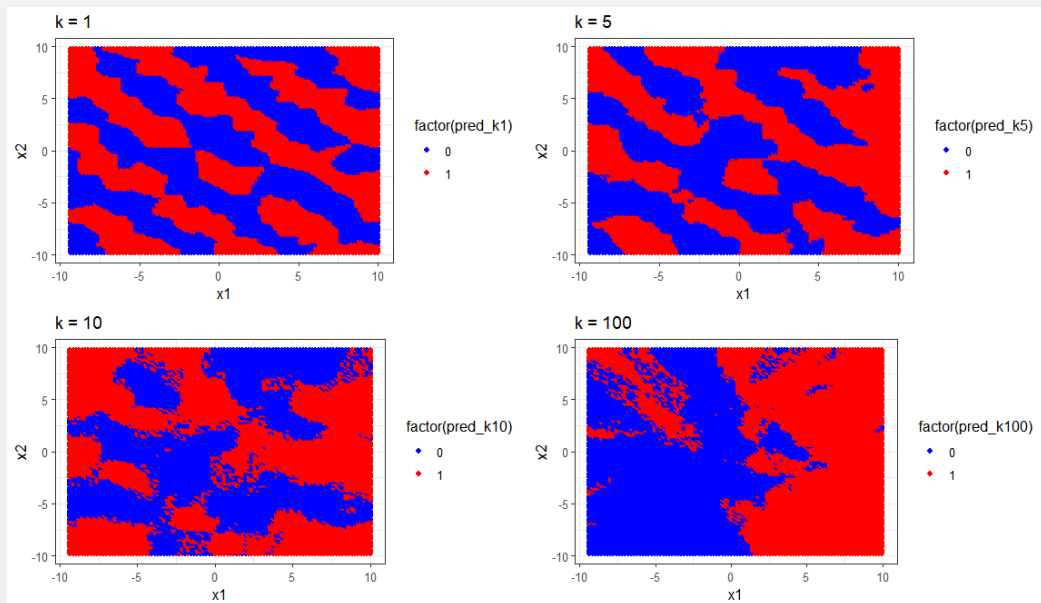
library(class)
k_values <- c(1, 5, 10, 100)
pred_list <- lapply(k_values, function(k) {
  knn(train = data[, c("x1", "x2")], test = grid, cl = data$y, k = k)
})
names(pred_list) <- paste0("k", k_values)

library(ggplot2)
grid$pred_k1 <- pred_list[["k1"]]
grid$pred_k5 <- pred_list[["k5"]]
grid$pred_k10 <- pred_list[["k10"]]
grid$pred_k100 <- pred_list[["k100"]]
p1 <- ggplot(grid, aes(x = x1, y = x2, color = factor(pred_k1))) +
  geom_point() +
  scale_color_manual(values = c("blue", "red")) +
  theme_bw() +
  ggtitle("k = 1")
p2 <- ggplot(grid, aes(x = x1, y = x2, color = factor(pred_k5))) +
  geom_point() +
  scale_color_manual(values = c("blue", "red")) +
```

```

theme_bw() +
  ggtitle("k = 5")
p3 <- ggplot(grid, aes(x = x1, y = x2, color = factor(pred_k10))) +
  geom_point() +
  scale_color_manual(values = c("blue", "red")) +
  theme_bw() +
  ggtitle("k = 10")
p4 <- ggplot(grid, aes(x = x1, y = x2, color = factor(pred_k100))) +
  geom_point() +
  scale_color_manual(values = c("blue", "red")) +
  theme_bw() +
  ggtitle("k = 100")
library(gridExtra)
grid.arrange(p1, p2, p3, p4, ncol = 2)

```



we can observe that,

For $k=1$, the decision boundary is very complex and follows the contours of the data.

As k increases, the decision boundary becomes smoother. For $k=10$ this approximates the true decision boundary much better.

$k=100$, the decision boundary becomes too simple and smooth. So, it will overfit.

- (2 points) For $k \in \{1, 2, 3, \dots, 100\}$ estimate the error rate of the model. At first, do so using the training error rate, i.e., the error the model has by predicting the original training data. At second, do so by using subsampling: Split the data into 2/3 training and 1/3 test data, train the model on the train data, test it on the test data. Repeat 100 times and use the mean value as an estimator for the error.

Solution:

```

n <- 250
x1 <- runif(n, -10, 10)
x2 <- runif(n, -10, 10)
y <- sin(x1 + x2) < 0

error_rate <- function(y_true, y_pred) { mean(y_true != y_pred) }
train_errors <- rep(NA, 100)

```

```

test_errors <- rep(NA, 100)

for (k in 1:100) {
  knn_model <- knn(train = cbind(x1, x2), test = cbind(x1, x2),
    cl = y, k = k)
  train_errors[k] <- error_rate(y, knn_model)
  test_errors_rep <- rep(NA, 100)
  for (i in 1:100) {
    set.seed(i)
    train_idx <- sample(1:n, size = round(2*n/3), replace = FALSE)
    test_idx <- setdiff(1:n, train_idx)

    knn_model <- knn(train = cbind(x1[train_idx], x2[train_idx]),
      test = cbind(x1[test_idx], x2[test_idx]),
      cl = y[train_idx], k = k)
    test_errors_rep[i] <- error_rate(y[test_idx], knn_model)
  }
  test_errors[k] <- mean(test_errors_rep)
}

```

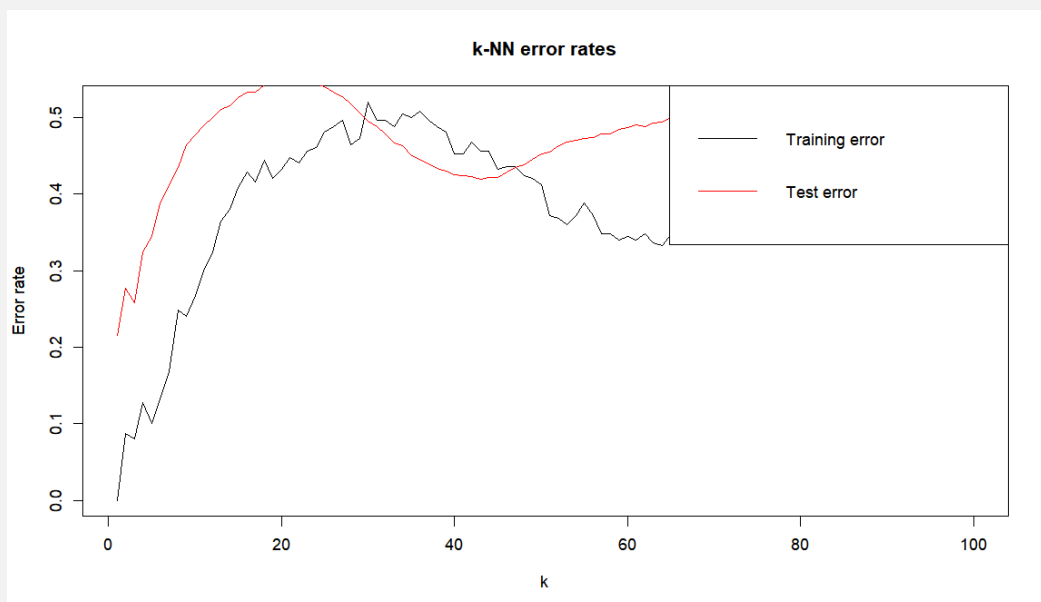
- Visualize both error rates with respect to k. What differences do you observe? How can you explain them?

Solution:

```

plot(1:100, train_errors, type = "l", xlab = "k",
  ylab = "Error rate", main = "k-NN error rates")
lines(1:100, test_errors, col = "red")
legend("topright", legend = c("Training error", "Test error"),
  lty = 1, col = c("black", "red"))

```



We can observe that the training error rate is always lower than the test error rate for all values of k. As we split the data into 2/3 training and 1/3 test data, it will perform better in training data.

And as the value of k increases, both the training and test error rates first increases upto a peak point and then tend to decrease with larger value of k .
Test error rates fluctuates in the middle range value of k .